

```
In [1]: import numpy as np
import pandas as pd
```

```
In [27]: df_tr = pd.read_excel('KPMG_VI_New_raw_data_update_final.xlsx', sheet_name='Transaction')
```

```
In [28]: df_tr['transaction_date'] = pd.to_datetime(df_tr['transaction_date'])
```

```
In [29]: df_tr['product_first_sold_date'] = pd.to_datetime(df_tr['product_first_sold_date'])
```

```
In [30]: df_tr.head()
```

```
Out[30]:
```

	transaction_id	product_id	customer_id	transaction_date	online_order	order_status	brand
0	1	2	2950	2017-02-25	0.0	Approved	Strider
1	2	3	3120	2017-05-21	1.0	Approved	Bicycle
2	3	37	402	2017-10-16	0.0	Approved	Oakley
3	4	88	3135	2017-08-31	0.0	Approved	NewBicycle
4	5	78	787	2017-10-01	1.0	Approved	GiantBicycle

```
In [9]: df_newcust = pd.read_excel('KPMG_VI_New_raw_data_update_final.xlsx', sheet_name='NewCustomerList')
```

```
/tmp/ipykernel_3439/450638472.py:1: FutureWarning: Inferring datetime64[ns] from data containing strings is deprecated and will be removed in a future version. To retain the old behavior explicitly pass Series(data, dtype=datetime64[ns])
df_newcust = pd.read_excel('KPMG_VI_New_raw_data_update_final.xlsx', sheet_name='NewCustomerList', skiprows=1)
```

```
In [10]: df_newcust.head()
```

```
Out[10]:
```

	first_name	last_name	gender	past_3_years_bike_related_purchases	DOB	job
0	Chickie	Brister	Male	86	1957-07-12	General Manager
1	Morly	Genery	Male	69	1970-03-22	Structural Engineer
2	Ardelis	Forrester	Female	10	1974-08-28	Senior Accountant
3	Lucine	Stutt	Female	64	1979-01-28	Account Representative
4	Melinda	Hadlee	Female	34	1965-09-21	Financial Analyst

5 rows × 7 columns

```
In [11]: df_custdemo = pd.read_excel('KPMG_VI_New_raw_data_update_final.xlsx', sheet_name='CustomerDemo')
```

```
/tmp/ipykernel_3439/2089153917.py:1: FutureWarning: Inferring datetime64
[ns] from data containing strings is deprecated and will be removed in a
future version. To retain the old behavior explicitly pass Series(data, d
type=datetime64[ns])
df_custdemo = pd.read_excel('KPMG_VI_New_raw_data_update_final.xlsx',sh
eet_name='CustomerDemographic',skiprows=1)
```

```
In [12]: df_custdemo.head()
```

```
Out[12]:
```

	customer_id	first_name	last_name	gender	past_3_years_bike_related_purchases	date_of_birth
0	1	Laraine	Medendorp	F	93	1953-10-01
1	2	Eli	Bockman	Male	81	1980-12-01
2	3	Arlin	Dearle	Male	61	1954-01-01
3	4	Talbot	NaN	Male	33	1961-10-01
4	5	Sheila-kathryn	Calton	Female	56	1977-05-01

```
In [13]: df_custaddr = pd.read_excel('KPMG_VI_New_raw_data_update_final.xlsx',sheet_name='CustomerAddress')
```

```
In [14]: df_custaddr.head()
```

```
Out[14]:
```

	customer_id	address	postcode	state	country	property_valuation
0	1	060 Morning Avenue	2016	New South Wales	Australia	10
1	2	6 Meadow Vale Court	2153	New South Wales	Australia	10
2	4	0 Holy Cross Court	4211	QLD	Australia	9
3	5	17979 Del Mar Point	2448	New South Wales	Australia	4
4	6	9 Oakridge Court	3216	VIC	Australia	9

Transactions DB

```
In [15]: df_tr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   transaction_id                        20000 non-null  int64
1   product_id                           20000 non-null  int64
2   customer_id                          20000 non-null  int64
3   transaction_date                     20000 non-null  datetime64[ns]
4   online_order                         19640 non-null  float64
5   order_status                         20000 non-null  object
6   brand                               19803 non-null  object
7   product_line                         19803 non-null  object
8   product_class                       19803 non-null  object
9   product_size                         19803 non-null  object
10  list_price                           20000 non-null  float64
11  standard_cost                       19803 non-null  float64
12  product_first_sold_date             19803 non-null  float64
dtypes: datetime64[ns](1), float64(4), int64(3), object(5)
memory usage: 2.0+ MB
```

```
In [16]: df_tr.isnull().sum()
```

```
Out[16]: transaction_id      0
product_id      0
customer_id     0
transaction_date 0
online_order    360
order_status    0
brand          197
product_line    197
product_class   197
product_size    197
list_price      0
standard_cost   197
product_first_sold_date 197
dtype: int64
```

- **Columns** `online_order`, `brand`, `product_line`, `product_class`, `product_size`, `standard_cost` **and** `product_first_sold_date` **have missing values. They can either be treated or dropped depending on the analysis**

```
In [18]: df_tr.nunique()
```

```
Out[18]: transaction_id      20000
product_id      101
customer_id     3494
transaction_date 364
online_order      2
order_status      2
brand             6
product_line      4
product_class     3
product_size      3
list_price       296
standard_cost     103
product_first_sold_date 100
dtype: int64
```

```
In [19]: df_tr.shape
```

```
Out[19]: (20000, 13)
```

- There are no duplicate transactions as the transaction_id column has all unique values

```
In [21]: df_tr.duplicated().sum()
```

```
Out[21]: 0
```

```
In [36]: df_tr['product_first_sold_date'].dt.date.unique()
```

```
Out[36]: array([datetime.date(1970, 1, 1), NaT], dtype=object)
```

- All values in the product_first_sold_date are from the same day, this can be an anomaly

New Customers DB

```
In [37]: df_newcust.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 23 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   first_name                           1000 non-null   object
 1   last_name                            971 non-null    object
 2   gender                               1000 non-null   object
 3   past_3_years_bike_related_purchases 1000 non-null   int64
 4   DOB                                  983 non-null    datetime64[ns]
 5   job_title                            894 non-null    object
 6   job_industry_category                835 non-null    object
 7   wealth_segment                       1000 non-null   object
 8   deceased_indicator                   1000 non-null   object
 9   owns_car                             1000 non-null   object
10   tenure                               1000 non-null   int64
11   address                              1000 non-null   object
12   postcode                             1000 non-null   int64
13   state                                1000 non-null   object
14   country                              1000 non-null   object
15   property_valuation                   1000 non-null   int64
16   Unnamed: 16                          1000 non-null   float64
17   Unnamed: 17                          1000 non-null   float64
18   Unnamed: 18                          1000 non-null   float64
19   Unnamed: 19                          1000 non-null   float64
20   Unnamed: 20                          1000 non-null   int64
21   Rank                                 1000 non-null   int64
22   Value                                1000 non-null   float64
dtypes: datetime64[ns](1), float64(5), int64(6), object(11)
memory usage: 179.8+ KB
```

- There are 5 columns unnamed, we'll have to drop them unless we can get the information from the company

```
In [52]: df_newcust.drop([i for i in df_newcust.columns if i.lower()[0:7] == 'unnamed'])
```

```
In [53]: df_newcust.columns
```

```
Out[53]: Index(['first_name', 'last_name', 'gender',  
               'past_3_years_bike_related_purchases', 'DOB', 'job_title',  
               'job_industry_category', 'wealth_segment', 'deceased_indicator',  
               'owns_car', 'tenure', 'address', 'postcode', 'state', 'country',  
               'property_valuation', 'Rank', 'Value'],  
              dtype='object')
```

```
In [54]: df_newcust.isnull().sum()
```

```
Out[54]: first_name          0  
         last_name          29  
         gender             0  
         past_3_years_bike_related_purchases  0  
         DOB                17  
         job_title          106  
         job_industry_category  165  
         wealth_segment      0  
         deceased_indicator    0  
         owns_car            0  
         tenure              0  
         address             0  
         postcode            0  
         state               0  
         country             0  
         property_valuation    0  
         Rank                0  
         Value               0  
         dtype: int64
```

- Columns `last_name`, `DOB`, `job_title` and `job_industry_category` have missing values. They can either be treated or dropped depending on the analysis

```
In [55]: df_newcust.duplicated().sum()
```

```
Out[55]: 0
```

- There are no duplicate values

Customer Demographic DB

```
In [56]: df_custdemo.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 13 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   customer_id                          4000 non-null   int64
 1   first_name                           4000 non-null   object
 2   last_name                            3875 non-null   object
 3   gender                               4000 non-null   object
 4   past_3_years_bike_related_purchases 4000 non-null   int64
 5   DOB                                  3913 non-null   datetime64[ns]
 6   job_title                            3494 non-null   object
 7   job_industry_category                3344 non-null   object
 8   wealth_segment                       4000 non-null   object
 9   deceased_indicator                   4000 non-null   object
10   default                              3698 non-null   object
11   owns_car                             4000 non-null   object
12   tenure                               3913 non-null   float64
dtypes: datetime64[ns](1), float64(1), int64(2), object(9)
memory usage: 406.4+ KB
```

```
In [57]: df_custdemo.isnull().sum()
```

```
Out[57]: customer_id          0
         first_name          0
         last_name        125
         gender            0
         past_3_years_bike_related_purchases  0
         DOB                87
         job_title         506
         job_industry_category  656
         wealth_segment      0
         deceased_indicator    0
         default           302
         owns_car            0
         tenure             87
dtype: int64
```

- **Columns** last_name , DOB , job_title , job_industry_category , default and tenure have missing values. They can either be treated or dropped based on the analysis

```
In [58]: df_custdemo.customer_id.nunique()
```

```
Out[58]: 4000
```

```
In [59]: df_custdemo.shape
```

```
Out[59]: (4000, 13)
```

```
In [60]: df_custdemo.duplicated().sum()
```

```
Out[60]: 0
```

- There are no duplicate values

```
In [61]: df_custdemo.gender.value_counts()
```

```
Out[61]: Female      2037
         Male       1872
         U          88
         F           1
         Femal       1
         M           1
         Name: gender, dtype: int64
```

- Values for gender haven't been captured properly or not cleaned

```
In [64]: df_custdemo.replace({'gender':{'M':'Male','Femal':'Female','F':'Female'}})
```

```
In [65]: df_custdemo.gender.value_counts()
```

```
Out[65]: Female      2039
         Male       1873
         U          88
         Name: gender, dtype: int64
```

- There are still 88 Unknown values

```
In [66]: df_custdemo.default.value_counts()
```

```
Out[66]: 100          113
         1           112
        -1          111
        -100         99
        Û;ÛçÛ£       53
        ...
        testâ testâ«  31
        /dev/null; touch /tmp/blns.fail ; echo  30
        âââtestââ    29
        ì,ëë°í ë¥´    27
        ,ââ»:*:â»ââ( â» Ì â» )ââ»:*:â»ââ    25
        Name: default, Length: 90, dtype: int64
```

- Default data has to be re-collected or dropped from the DB

Customer address DB

```
In [67]: df_custaddr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3999 entries, 0 to 3998
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   customer_id           3999 non-null   int64
 1   address                3999 non-null   object
 2   postcode               3999 non-null   int64
 3   state                  3999 non-null   object
 4   country                3999 non-null   object
 5   property_valuation     3999 non-null   int64
dtypes: int64(3), object(3)
memory usage: 187.6+ KB
```

```
In [70]: df_custaddr.isnull().sum()
```

```
Out[70]: customer_id      0
         address          0
         postcode         0
         state            0
         country          0
         property_valuation 0
         dtype: int64
```

- No missing values

```
In [71]: df_custaddr.duplicated().sum()
```

```
Out[71]: 0
```

```
In [72]: df_custaddr.customer_id.nunique()
```

```
Out[72]: 3999
```

```
In [73]: df_custaddr.shape
```

```
Out[73]: (3999, 6)
```